

Poisson Regression

Estimation, Interpretations and Diagnostics

1

Last week: Logistic Regression

- Logistic regression: definition *for binary responses.*
- MLE Estimation and Interpretation for different type of covariates
- about β_1 Inference: confidence intervals and tests
- Fitted values and residuals
- Diagnostics: overdispersion

→ similar to linear regression models except that the covariate effects are on log(odds) of the response, rather than the response itself.

$$\rightarrow E(y) = P(y=1) = \frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}}$$

$$Var(y) = E(y)(1 - E(y))$$

under Bernoulli dist. assumption.

2

This week: Poisson Regression

Poisson regression is a regression analysis designed for modelling count data.

The response variable represents the number of times an event occurs in a given timeframe or even in a space such as a geographic unit. For example,

- the number of customer arrivals
- disease occurrences
- traffic accidents

Poisson dist. $P(Y=k) = \frac{e^{-\lambda} \times \lambda^k}{k!}$,
 $k=0, 1, 2, 3, \dots$

$$E(Y) = \lambda$$
$$Var(Y) = \lambda$$

The response

Scroll down to see full content

Poisson is a common distribution to model a random variable that takes discrete non-negative integer values (i.e., 0, 1, 2,...)

$$Y_i | \mathbf{X}_i \sim \text{Poisson}(\lambda_i),$$

where λ_i is the conditional mean.

A particularity of the Poisson distribution is that its mean equals its variance:

$$\lambda_i = E[Y_i | \mathbf{X}_i] = Var[Y_i | \mathbf{X}_i]$$

Thus, any factor that affects the mean will also affect the variance.

This fact could be a potential drawback for using a Poisson regression model.

Poisson regression can also be used to model rates of events but we won't cover that case in this course.

linear model
 $Var(Y_i | X_i) = \sigma^2$
(indep. of β_s)

Poisson model
 $E(Y_i | X_i) = \beta_0 + \beta_1 X_i$
 $Var(Y_i | X_i) = \beta_0 + \beta_1 X_i$

Dataset

In this lecture, we'll use the [Bikeshare dataset](#) from the [ISL](#) package discussed in the [ISL book](#)

The dataset contains information on the hourly users of a bike sharing program in Washington DC, for different months, hours, weather conditions, and membership types (see ? [Bikeshare](#))

The Bikeshare dataset has information about 8645 users

season	mnth	day	hr	holiday	weekday	workingday	weathersit	temp	atemp
3	Aug	220	6	0	1	1	clear	0.70	0.6
4	Oct	293	21	0	4	1	clear	0.44	0.4

5

A range problem

Since counts are nonnegative integers, if we use linear regression, we might have a range problem by predicting negative counts.

Instead, we can use a *function* of the linear component with a positive range, for example, an exponential function.

this is not the only function but it's commonly used due to its simplicity and interpretation

6

Poisson Regression

Scroll down to see full content

The Poisson regression is defined as:

$$E[Y_i | \mathbf{X}_i] = \lambda_i = e^{\beta_0 + \beta_1 X_{1,i} + \dots + \beta_p X_{p,i}} \quad [\text{mean}]$$

or equivalently,

$$\log(\lambda_i) = \beta_0 + \beta_1 X_{1,i} + \dots + \beta_p X_{p,i} \quad [\text{log-mean}]$$

the log function links the conditional expectation with the linear component in a natural way in the GLM family so it's called the canonical link.

We model the log of the mean counts with a linear component!!

then all that we learned for MLR can be used for this function of the conditional expectation!!

Note that there is no error term! we are modeling a function of the conditional expectation

Y_i counts

linear predictor L_i

not a r.v., so no e

$E(Y_i)$

$g(E(Y_i | X_i)) = L_i$
where $g(y) = \log(y)$
is a link function

in a logistic model,
 $g(y) = \log \frac{y}{1-y}$
in a linear model
 $g(y) = y$

Poisson: a generalized linear model

The Poisson regression also belongs to the generalized linear models (GLMs) family.

In R, we can use the `glm()` function again but setting `family = poisson` to calculate the *maximum likelihood estimator* of the coefficients.

`glm()` models, by default, the log of the mean counts, which is the "canonical link" function but other functions can be used

Fitting a Poisson Regression

Scroll down to see full content

The interpretation of the coefficients will be very similar to those of the logistic regression case, *except that the covariate effects are on log-scale instead of log(odds) scale.*

So this time let's fit a model with various types of covariates at once.

From all the covariates available, let's consider

X₁ • **workingday**: an indicator variable that equals 1 if it is neither a weekend nor a holiday

X₂ • **temp**: the normalized temperature, in Celsius

X₃ • **windspeed**: the normalized windspeed

X₄ • **season**: a qualitative variable that takes four values for each season
categorical variable with 4 levels.

Factors vs numeric

Note that R will create dummy variables of factors, not of numerical variables.

sometimes numerical values are used for qualitative variables, e.g., season.

```
# A tibble: 2 × 5
  season workingday temp windspeed bikers
  <dbl>      <dbl> <dbl>    <dbl>   <dbl>
1     1         0  0.26    0.134    34
2     1         1  0.28    0.104    96
```

You will need to adjust that before fitting the model

Let's fit a Poisson model to the data as is. → include all predictors

```
1 summary(glm(formula = bikers ~ ., data = bikeshare,
2             family = 'poisson'))
```

Call:
glm(formula = bikers ~ ., family = "poisson", data = bikeshare)

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	3.411649	0.004158	820.39	<2e-16 ***
season	0.123221	0.000969	127.16	<2e-16 ***
workingday	-0.028971	0.001940	-14.93	<2e-16 ***
temp	2.050466	0.004919	416.82	<2e-16 ***
windspeed	0.822224	0.007361	111.70	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 1052921 on 8644 degrees of freedom
Residual deviance: 805356 on 8640 degrees of freedom
AIC: 858404

← X should have 3 estimates for "season"

← ? If this is not true, the above results are wrong!

Note that **workingday** and **season** were fit as numerical variables, and both should be treated as categorical variables!! ✓

11

R can't distinguish continuous from categorical variables if the variable is numeric.

Factors are required for qualitative variables!

```
1 bikeshare <- bikeshare %>%
2   mutate(season = as.factor(season),
3          workingday = as.factor(workingday))
```

Call:
glm(formula = bikers ~ ., family = "poisson", data = bikeshare)

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	3.345358	0.003981	840.38	<2e-16 ***
season2	0.062916	0.003741	16.82	<2e-16 ***
season3	-0.141267	0.004320	-32.70	<2e-16 ***
season4	0.379421	0.003357	113.01	<2e-16 ***
workingday1	-0.036254	0.001942	-18.67	<2e-16 ***
temp	2.688047	0.007435	361.54	<2e-16 ***
windspeed	0.741449	0.007377	100.50	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

log-mean vs mean counts

As before, we can interpret the “raw” results corresponding to log-mean counts and are interpreted as for MLR

$$\log(E[\text{bikers}|\mathbf{X}]) = \beta_0 + \beta_1 S_2 + \beta_2 S_3 + \beta_3 S_4 + \beta_4 W_1 + \beta_5 \text{temp} + \beta_6 \text{wind}$$

for simplicity, S = season, W = workingday, and **wind** = windspeed, and $\mathbf{X}$ = all covariates

three dummy variables for "season"
the reference category is hidden

$S_2 = \begin{cases} 1 & \text{season 2} \\ 0 & \text{otherwise} \end{cases}$
 $S_3 = \begin{cases} 1 & \text{season 3} \\ 0 & \text{otherwise} \end{cases}$

13

Interpretation: log-mean

Scroll down to see full content

term	estimate	std.error	statistic	p.value
(Intercept)	3.345	0.004	840.384	0
season2	0.063	0.004	16.818	0
season3	-0.141	0.004	-32.699	0
season4	0.379	0.003	113.008	0
workingday1	-0.036	0.002	-18.672	0
temp	2.688	0.007	361.541	0
windspeed	0.741	0.007	100.505	0

iClicker: true or false

- Keeping other variables in the model constant, at any value, the mean bike usage decrease 0.036 in working days

A: true B: false

- Keeping other variables in the model constant, at any value, a unit increase in the normalized temperature is associated with an increase in log mean bike usage of 2.688.

A: true B: false

14

Interpretation: log-mean

The coefficients can be interpreted in terms of the log-means:

- **dummy variables** $\hat{\beta}_1 = 0.063$: the log-mean usage is season 2 is 0.063 larger than that in season 1 (reference), keeping all other variables in the model constant at any value
- **continuous variables** $\hat{\beta}_5 = 2.688$: an increase of 1 unit in **temp** is associated with an increase in the log-mean usage of 2.688, keeping all other variables in the model constant at any value

same as for MLR but with response “log-mean counts”!

15

Keeping constant ...

$$\log \lambda_{\text{temp}} = \beta_0 + \beta_5 \widehat{T}^{\text{temp}} + \overbrace{\beta_1 S2 + \beta_2 S3 + \beta_3 S4 + \beta_4 W1 + \beta_6 \text{wind}}^{\text{Constant}}$$

$$\log \lambda_{\text{temp} + 1} = \beta_0 + \beta_5 \widehat{T}^{\text{temp}+1} + \overbrace{\beta_1 S2 + \beta_2 S3 + \beta_3 S4 + \beta_4 W1 + \beta_6 \text{wind}}^{\text{Constant}}$$

We take the difference between both equations as:

$$\begin{aligned} \log \lambda_{\text{temp} + 1} - \log \lambda_{\text{temp}} &= \beta_5 (T + 1) - \beta_5 T \\ &= \beta_5 \end{aligned}$$

Thus, we interpret these coefficients as **difference in log-mean counts**

16

Interpretation: mean counts

Scroll down to see full content

But by logarithm property for a ratio:

$$\log \frac{\lambda_{\text{temp} + 1}}{\lambda_{\text{temp}}} = \log \lambda_{\text{temp} + 1} - \log \lambda_{\text{temp}} \\ = \beta_5$$

Finally, we have to exponentiate the previous equation:

$$\frac{\lambda_{\text{temp} + 1}}{\lambda_{\text{temp}}} = e^{\beta_5}$$

This expression indicates that an increase of 1 unit in a continuous covariate is associated with a mean count change **by a factor of e^{β_5}** (multiplicative change).

$$\lambda_{\text{temp} + 1} = e^{\beta_5} \lambda_{\text{temp}}$$

A more natural way of interpreting the coefficients of a Poisson regression

Poisson model:

$$\log(E(y)) = \beta_0 + \beta_1 x_1 + \beta_2 x_2$$

$$E(y) = e^{\beta_0} \times e^{\beta_1 x_1} \times e^{\beta_2 x_2}$$

Using tidy: exponentiate = TRUE

```
1 #Coefficients for the mean counts
2 tidy(bikers_poisson, exponentiate = TRUE) %>%
3   mutate_if(is.numeric, round, 3) %>%
4   knitr::kable()
```

term	estimate	std.error	statistic	p.value
(Intercept)	28.371	0.004	840.384	0
season2	1.065	0.004	16.818	0
season3	0.868	0.004	-32.699	0
season4	1.461	0.003	113.008	0
workingday1	0.964	0.002	-18.672	0
temp	14.703	0.007	361.541	0
windspeed	2.099	0.007	100.505	0

e^{β_j}

18

Scroll down to see full content

- **dummy variables** $e^{\hat{\beta}_1} = 1.065$: the mean usage is season 2 increases by a factor of 1.065 that in season 1 (reference), keeping all other variables in the model constant at any value
 - or by $(1.065 - 1) \times 100\% = 6.5\%$
- **continuous variables** $e^{\hat{\beta}_5} = 14.703$: an increase of 1 unit in **temp** is associated with an increase in mean usage of a factor of 14.703, keeping all other variables in the model constant at any value

iClicker:

dummy variables S3 $e^{\hat{\beta}_2} = -0.141 = 0.868$:

A: a change in season from season 1 to season 3 is associated with a change in mean bike usage by a factor of 0.868, keeping all other variables in the model constant at any value

B: only 86.8% as many people will use bikes in season 3 relative to season season 1, keeping all other variables in the model constant at any value

C: a change in season from season 1 to season 3 is associated

19

Again, 4 Poisson regressions in 1!! For the **additive** model:

$$\log \lambda = \beta_0 + \beta_1 S_2 + \beta_2 S_3 + \beta_3 S_4 + \beta_4 W_1 + \beta_5 \text{temp} + \beta_6 \text{wind}$$

$$\log \lambda = \begin{cases} \beta_0 + \beta_4 W_1 + \beta_5 \text{temp} + \beta_6 \text{wind} & \text{if } S_2 = S_3 = S_4 = 0 \\ (\beta_0 + \beta_1) + \beta_4 W_1 + \beta_5 \text{temp} + \beta_6 \text{wind} & \text{if } S_2 = 1, S_3 = S_4 = 0 \\ (\beta_0 + \beta_2) + \beta_4 W_1 + \beta_5 \text{temp} + \beta_6 \text{wind} & \text{if } S_3 = 1, S_2 = S_4 = 0 \\ (\beta_0 + \beta_3) + \beta_4 W_1 + \beta_5 \text{temp} + \beta_6 \text{wind} & \text{if } S_4 = 1, S_2 = S_3 = 0 \end{cases}$$

reference season

dummy variables

(optional but useful) try to write these in exponential form

20

Model with interaction

For simplicity, let's look at a model with only **workingday** and **temp** but interacting

we are thinking that the "effect" of temperature will be different for leisure days compared to working days.

$$\log \lambda = \gamma_0 + \gamma_1 W_1 + \gamma_2 \text{temp} + \gamma_3 W_1 \times \text{temp}$$

21

Again, 2 Poisson regressions in 1 with **different** slopes!

Again, 2 Poisson regressions in 1 with **different** slopes!

$$\log \lambda = \gamma_0 + \gamma_1 W1 + \gamma_2 \text{temp} + \gamma_3 W1 \times \text{temp}$$

For the **log-mean counts**:

$$\log \lambda = \begin{cases} \gamma_0 + \gamma_2 \text{temp} & \text{if } W1 = 0 \\ (\gamma_0 + \gamma_1) + (\gamma_2 + \gamma_3) \text{temp} & \text{if } W1 = 1 \end{cases}$$

different *slope* (γ_2 vs $(\gamma_2 + \gamma_3)$), different *intercepts* (γ_0 vs $(\gamma_0 + \gamma_1)$)!

For the **mean counts**:

$$\lambda = \begin{cases} e^{\gamma_0} e^{\gamma_2 \text{temp}} & \text{if } W1 = 0 \\ e^{(\gamma_0 + \gamma_1)} e^{(\gamma_2 + \gamma_3) \text{temp}} & \text{if } W1 = 1 \end{cases}$$

22

in R

Scroll down to see full content

```
1 bikers_poisson_int <- glm(bikers~workingday*temp,  
2                           data = bikeshare,  
3                           family = poisson)  
4  
5 tidy(bikers_poisson_int, exponentiate = TRUE) %>%  
6   mutate_if(is.numeric, round, 3) %>%  
7   knitr::kable()
```

term	estimate	std.error	statistic	p.value
(Intercept)	34.343	0.005	687.857	0
workingday1	1.481	0.006	63.450	0
temp	14.723	0.008	318.215	0
workingday1:temp	0.483	0.010	-71.201	0

include interaction

significant interaction

23

Interpretation (for mean counts)

Scroll down to see full content

In a **model with interactions**, the expected change in one variable **depends on** the value or levels the other variable

We **CAN'T** say “keeping the **temp** value constant” or “for working days or leisure days”

- for **leisure days**, a 1 unit increase in temperature is associated with an increase of mean bike usage **by a factor of** 14.723.

24

Interaction term

To get the “effect” of temperature for working days (i.e., **workingday** = 1), you need to *multiply* (instead of summing) the estimated coefficient of **temp** by that of the interaction term:

From the **tidy** table, with **exponentiate** = TRUE.

$$e^{\hat{\gamma}_2} = 14.723, e^{\hat{\gamma}_3} = 0.483$$

We need to multiply (instead of summing):

$$e^{\hat{\gamma}_2 + \hat{\gamma}_3} = e^{\hat{\beta}_2} e^{\hat{\beta}_3} = 14.723 \times 0.483 = 7.11$$

In **working days**, a unit increase in temperature is associated with an increase of mean usage **by a factor of** 7.11.

25

Inference

Inference for Poisson regression is the same as that of Logistic regression. They are based on theoretical results from the CLT.

if the sample size is large, the *Normal* distribution is used to obtain boundaries of confidence intervals and *p*-values

$$\frac{\hat{\beta}_j}{\hat{SE}(\hat{\beta}_j)} \sim N(0, 1)$$

CLT

This is called the **Wald's statistic**

26

Scroll down to see full content

The confidence intervals can be calculated for both the *log-mean* and the *mean* counts.

But again, R will do all this for you!

```
1 tidy(bikers_poisson,  
2   conf.int = TRUE, conf.level = 0.9) %>%  
3   mutate_if(is.numeric, round, 3) %>%  
4   knitr::kable()
```

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	3.345	0.004	840.384	0	3.339	3.352
season2	0.063	0.004	16.818	0	0.057	0.069
season3	-0.141	0.004	-32.699	0	-0.148	-0.134
season4	0.379	0.003	113.008	0	0.374	0.385
workingday1	-0.036	0.002	-18.672	0	-0.039	-0.033
temp	2.688	0.007	361.541	0	2.676	2.700
windspeed	0.741	0.007	100.505	0	0.729	0.754

> options(digits=2)

reject Ho

all exclude 0

wrong

Residuals

The Poisson regression has the same problems as the Logistic Regression:

- responses are discrete (at least more than 2 values now)
- the variance is not constant

don't use this one

$$\text{resid}_{\text{raw}} = y_i - \hat{\lambda}_i,$$


use this one

$$\text{resid}_{\text{pearson}} = \frac{y_i - \hat{\lambda}_i}{\sqrt{\hat{\lambda}_i}} \quad \text{SE} \quad \checkmark$$

28

Residual plot

```
1 plot(bikers_poisson,1)
```



29

Overdispersion

The variance of a Poisson random variable equals the mean, which in practice may not be observed.

Poisson regression usually exhibits overdispersion.

An easy way to check is to fit a model using the quasilikelihood method and check if the dispersion parameter is very different from 1.

Poisson model

$$\begin{cases} E(y) = \beta_0 + \beta_1 x \\ \text{Var}(y) = E(y) \end{cases}$$

here $\delta = 1$

$$\begin{cases} E(y) = \beta_0 + \beta_1 x \\ \text{Var}(y) = \delta \times E(y) \end{cases}$$

called dispersion parameter

30

in R

```
1 bikers_quasipoisson <- glm(bikers ~ .,
2                             data = bikeshare,
3                             family = quasipoisson)
4
5 summary(bikers_quasipoisson)
```

Call:
glm(formula = bikers ~ ., family = quasipoisson, data = bikeshare)

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.34536	0.03789	88.286	< 2e-16 ***
season2	0.06292	0.03561	1.767	0.077302 .
season3	-0.14127	0.04112	-3.435	0.000595 ***
season4	0.37942	0.03196	11.872	< 2e-16 ***
workingday1	-0.03625	0.01848	-1.962	0.049849 *
temp	2.68805	0.07077	37.981	< 2e-16 ***
windspeed	0.74145	0.07022	10.558	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for quasipoisson family taken to be 90.60984)

Null deviance: 1852921 on 8644 degrees of freedom
Residual deviance: 783300 on 8638 degrees of freedom
AIC: NA

Number of Fisher Scoring iterations: 5

drop1(fit1)

more reliable than previous results.

very far from 1, so there is a serious overdispersion problem, i.e. Poisson dist does not hold.

$$\delta = 90.609$$

Summary

Scroll down to see full content

- A count response can be modelled using a Poisson Regression
- The (conditional) expectation of a count response is the (conditional) *mean counts*.
- The (conditional) expectation of a count response equals the (conditional) variance
- A LR can not be used to model the conditional expectation of a count response since its range extends beyond positive values
- Instead, one can model a function of the conditional mean counts. A common choice in Poisson regression is to use the *log* function

32

Summary (cont.)

The interpretation of the coefficients depends on the type of variables and the form of the model:

The raw coefficients (of the log-mean count model) are interpreted as:

- log-mean counts of a reference group
- difference of log-mean counts of a treatment vs a control group
- association between unit change in the input and changes in log-mean counts

33

Summary (cont.)

The **exponentiated coefficients** (of the mean count model) are interpreted as:

- mean counts of a reference group
- ratio of mean counts of a treatment vs a control group
- multiplicative changes in mean counts per unit change in the input

34

Summary (cont.)

- The estimated Poisson model can be used to make inference using the Wald's statistic: confidence intervals and tests
- There are different type of fitted values (log-mean counts or counts?) and residuals
- ✓ Overdispersion is a potential problem in Poisson regression. Using *quasilikelihood* approach usual helps.

35