

Model Evaluation

Goodness of Fit

1

Part 1: Model Estimation and Tests

- Linear regression: continuous response
- Logistic regression: binary response
- Poisson regression: count response
- Diagnostics of assumptions

iterative algorithms

Common issues:
1. multicollinearity
2. interaction

may not converge

GLM

$$g(E(y)) = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p$$

where $g(\cdot)$ is a link function.
 ↳ smooth, monotone function

linear predictor

This week: Model Evaluation



Photo by Siora Photography on Unsplash

Does the model fit the data well? Do we need more/less variables in the model?

3

This week

We will evaluate different linear models estimated by least squares (LS)

Today: Goodness of fit

- the coefficient of determination, R^2
- other evaluation metrics

Next class: comparing nested models

- F -test to compare our model against an intercept-only model
- a more general F -test to compare nested models

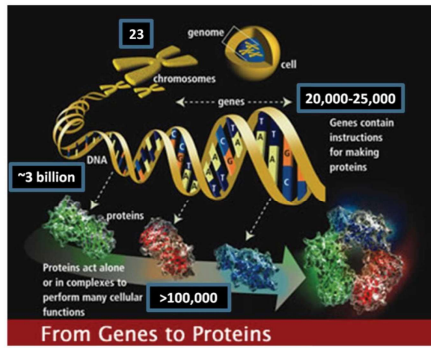
100% 20%-50%
percent of variation in the response that can be explained by the model
 $y = \beta_0 + \beta_1 x + e$
noise
noise
EQ
objective: where does the variation in y come from?

Case Study: protein vs mRNA

We will work with a new case study from Biology aiming to study the relation between mRNA and protein levels

- the Central Dogma of Biology: protein levels should be highly related to mRNA levels
- however, most experimental data found weak associations between these features

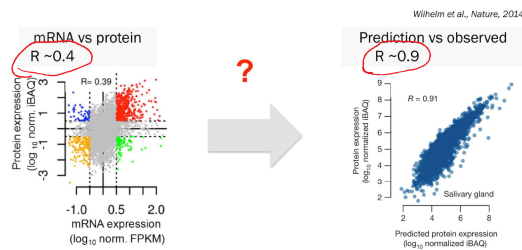
Can we predict protein from mRNA levels??



http://www.ornl.gov/sci/techresources/Human_Genome/project/info.shtml

6

An ambitious result



7

An ambitious claim

Scroll down to see full content

In 2014, Wilhelm et al. proposed a method to predict protein from mRNA levels and claimed:

[...] it now becomes possible to predict protein abundance in any given tissue with good accuracy from the measured mRNA

For each gene g the predicted protein level in the t -th tissue is given by

$$\hat{\text{prot}}_t = \hat{r}_g * \text{mrna}_t$$

where, $\hat{r}_g$ is the median of all the prot/mrna ratios of gene g

The subscript t is to identify each tissue.

The subscript g is to emphasize that the slope is *gene-specific*.

Models (examples for 3 genes)

Scroll down to see full content

- model.1:

null model

$$\text{prot}_t = \beta_0 + \varepsilon_t$$

treatment effect on ^y reduction

null model, intercept only model
(none of the predictors are useful)

- model.2:

$$\text{prot}_t = \beta_0 + \beta_1 \text{mrna}_t + \varepsilon_t$$

nested

- model.3:

$$\text{prot}_t = \beta_0 + \beta_2 \text{gene2}_t + \beta_3 \text{gene3}_t + \varepsilon_t$$

nested

- model.4:

$$\text{prot}_t = \beta_0 + \beta_1 \text{mrna}_t + \beta_2 \text{gene2}_t + \beta_3 \text{gene3}_t + \varepsilon_t$$

nested

In R, stepwise variable
selection:

full
> stepAIC(model)

forward
backward

reasonable

→ start with null model
→ add x_1 , check R^2
→ add x_2 , check R^2
...

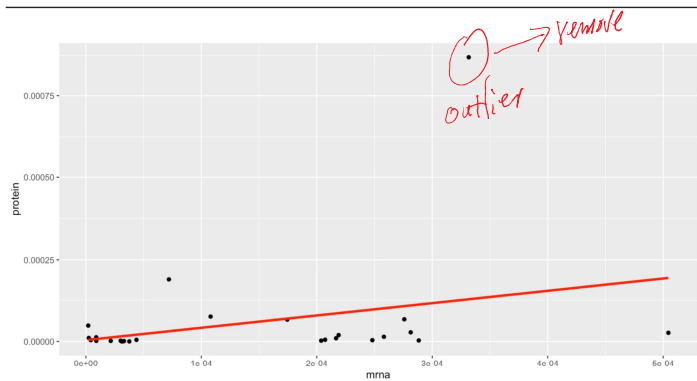
model.2: SLR

```
1 set.seed(561)
2 dat_3genes <- dat_bio %>%
3   subset(gene %in% sample(gene,3))
4
5 model2 <- lm(prot~mrna, data=dat_3genes)
6 tidy(model2)
```

A tibble: 2 × 5

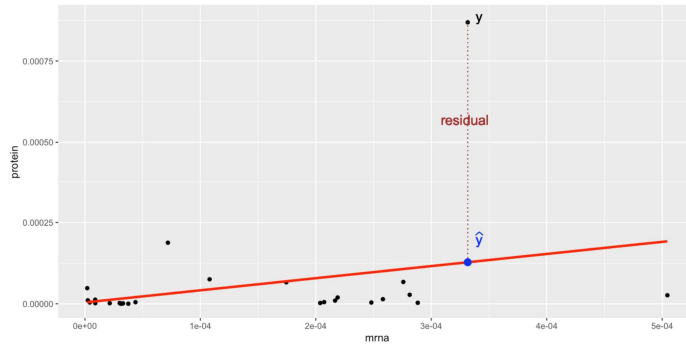
term	estimate	std.error	statistic	p.value
<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1 (Intercept)	0.00000409	0.0000475	0.0060	0.932
2 mrna	0.373	0.247	1.51	0.144

10



11

Prediction and residuals



12

The predicted values

The blue dot, $\hat{y}$ in the line, is the *predicted* protein level for a given amount of mRNA in this gene.

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$$

we put a “hat” on y to indicate that it’s a predicted value using the estimated regression

NOTE: there’s no error term in the predicted model!

13

Predicted values

Scroll down to see full content

- They are random variables since they are functions of the estimated coefficients

we can estimate their standard error and CI!! (coming next class!)

- In R output all predicted values are stored in the column `.fitted`.

In our case:

$$\hat{\text{prot}}_t = 4.18 \times 10^{-6} + 0.37 \text{ mrna}_t$$

The residuals

Scroll down to see full content

The *residual* is the difference between the predicted and the observed value of the response

In R output all residuals are stored in the column `.resid`.

Note that:

$$\text{res}_i = y_i - \hat{y}_i$$

observed → y_i $\hat{y}_i$ → *predicted*

$\varepsilon_i \neq \text{res}_i$

The residuals are the prediction errors.

In our case:

$$\text{res}_t = \text{prot}_t - \hat{\text{prot}}_t$$

Goodness of fit

Is our model better than “nothing”?

Scroll down to see full content

→ null model → $\hat{y}_i = \bar{y}$

- The best predictor of the response Y is $E[Y]$ which we can estimate with the sample mean of Y ...
- Given the (additional) information in $\mathbf{X}$, the best predictor is $E[Y|\mathbf{X}]$

“nothing” means no explanatory variables, intercept-only model, aka null model

So, an important question is:

Is our linear regression better than just using $E[Y]$ to predict??

Statistically, we want to compare our prediction $\hat{Y}$ (an estimate of $E[Y|\mathbf{X}]$) with $\bar{Y}$ (an estimate of $E[Y]$)

Comparing predictions

Explained Sum of Squares: $ESS = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2$

- $\hat{y}_i$ predicts y_i using the LR

- $\bar{y}$ predicts y_i without a model.

If our model is better than nothing, this should be large!!

The ESS measures how much variation in the data is explained by the additional information given by the LR model

variation explained by your model

SS measures variation

$\hat{y}_i$ is the predicted value based on a model
 $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$

$\bar{y}$ is the predicted value based on the null model

Residual Sum of Squares: $RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$

- this is the sum of the squares of the residuals from the fitted model

- our estimated parameters minimize these distances!!

variation that cannot be explained by the model (or variation from noise)

residuals

$$y = \underbrace{\beta_0 + \beta_1 x}_{\text{model}} + \underbrace{e}_{\text{error}}$$

Total Sum of Squares: $TSS = \sum_{i=1}^n (y_i - \bar{y})^2$ total variation

- this is the sum of the squares of the residuals from the null (intercept-only, no explanatory variables) model
- when properly scaled, it is the sample variance of Y which estimates the population variance of Y

Q: How much is the total variation in the data (response data) that can be explained by your model?

Sample variance $s^2 = \frac{1}{n-1} \underbrace{\sum_{i=1}^n (y_i - \bar{y})^2}_{TSS}$

Sum of squares decomposition

If parameters are estimated using LS and the LR has an intercept, $TSS = ESS + RSS$

$$\sum_{i=1}^n (y_i - \bar{y})^2 = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2 + \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$TSS = ESS + RSS$

$N(y_i)^2 \rightarrow x^2$
decomposition of total variation

ANOVA: similar idea

Q: How to show that your model is more useful than the null model?

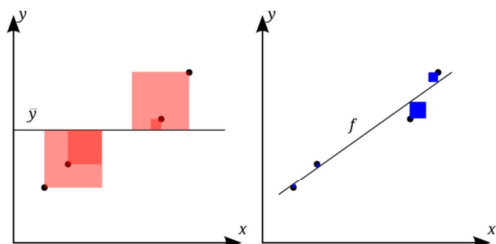
A: Compare ESS to RSS (or TSS) to see if ESS is much larger.

The coefficient of determination

[Scroll down to see full content](#)

If our model provides a good fit, we expect the TSS (residuals from the null model, in red) to be much larger than the RSS (residuals from the fitted model, which we minimized by LS, in blue)!!

Picture from Wikipedia



Using the decomposition above and dividing by TSS:

$$1 = \frac{ESS}{TSS} + \frac{RSS}{TSS}$$

The Coefficient of determination was first defined as

$R^2 = \frac{ESS}{TSS} = 1 - \frac{RSS}{TSS}$

= proportion of total variation that can be explained by the model

Interpretations

For a LR with an intercept and estimated by LS, the coefficient of determination: R^2

- measures the gain in predicting the response using the LM instead of the response sample mean, relative to the total variation in the response.
- is also interpreted as the proportion of ^{total} variance of the response (TSS) explained by the model (ESS)
- is between 0 and 1 since we expect TSS to be much larger than RSS (thus their ratio is smaller than 1)

Statistical Thinking

$$0 \leq R^2 = 1 - \frac{RSS}{TSS} \leq 1$$

22

1.3 Computing in R

Don't worry, R computes this statistic for you!!

For the full model

```
1 glance(model2)
# A tibble: 1 x 12
#   r.squared adj.r.squared sigma statistic p.value    df logLik   AIC    BIC
#   <dbl>      <dbl>      <dbl>    <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>
1  0.0869      0.0489  0.000166    2.29    0.144     1  190. -375. -371.
#   3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

Choose a model with a small AIC value.

AIC = (goodness-of-fit) + ~~penalty~~ for large model (or large # of predictors)

BIC similar

23

$R^2 = 0.08$
↑
too small

- For the 3 genes randomly selected in this example, only 9% of the total variation in protein levels is explained by *mrna*!!
- The coefficient of determination equals the square of the coefficient of multiple correlation

you probably noticed this in the SLR: $r^2 = R^2$

in fact, the coefficient of determination was first introduced this way in 1921!

24

For *gene-specific* models

Let's look at each gene separately

```

      gene values.av tissue      prot      mrna
1308 ENSG00000110075      7 uterus      NA 0.0000337
1324 ENSG00000110700     12 uterus 2.767271e-05 0.0002813
2572 ENSG00000139194      7 uterus 1.871157e-06 0.0000090
6132 ENSG00000110075      7 kidney      NA 0.0000269
6148 ENSG00000110700     12 kidney 6.606698e-05 0.0001744
7396 ENSG00000139194      7 kidney 8.691741e-04 0.0003316

```

```

1 dat_3genes %>% group_by(gene) %>%
2   do(glance(lm(prot ~ mrna, data = .)))

```

```

# A tibble: 3 × 13
# Groups:   gene [3]
  gene      r.squared adj.r.squared   sigma statistic p.value    df logLik
  <fct>      <dbl>      <dbl>    <dbl>      <dbl>  <dbl>  <dbl> <dbl>
1 ENSG0000... 0.179      0.0142 1.69e-6    1.09 3.45e-1    1 84.3
-163.
2 ENSG0000... 0.0609     -0.0330 2.77e-5    0.648 4.39e-1    1 110.
-214.

```

25

Is the R^2 useful??

Scroll down to see full content

When we use LS, we get an $R^2 = 0.56$, is this value large??

- The R^2 can be used to compare the size of the residuals of the fitted model with those of the null
- The R^2 can't be used to test any hypothesis to answer this question since its distribution is unknown
- The R^2 is computed based on *in-sample* observations
- The R^2 does not provide a sense of how good is our model in predicting *out-of-sample* cases (aka test set)!!
- Note that R^2 computed as $R^2 = 1 - \frac{RSS}{TSS}$ ranges between 0 and 1 if the LR model has an intercept and is estimated by LS!!
- ✓ The R^2 increases as new variables are added to the model, regardless of their relevance!! Thus, it can't be used to compare nested models

56% of total variation in the response data can be explained by the model

R^2 increases as the # of predictors increase in the model
⇒ large models

Model Evaluation via Adjusted R^2

Scroll down to see full content

The R^2 increases as more input variables are added to the model since the $RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$ decreases as more input variables are included in the model.

To overcome this issue with R^2 , we can obtain an **adjusted R^2** as follows:

$$\text{adjusted } R^2 = 1 - \frac{RSS/(n - p - 1)}{TSS/(n - 1)},$$

where

- p is the number of covariates in the model (excluding the intercept)
- n is our sample size to estimate the model

This adjusted coefficient of determination penalizes RSS with

penalty for large p .

usefulness of your model in linear models

Other evaluation metrics

Scroll down to see full content

Other related metrics used to evaluate LR which measure how far Y is from $\hat{Y}$ are:

Residuals Standard Error: $RSE = \sqrt{\frac{1}{n-p-1} RSS}$

called `sigma` in `glance()` output

→ Variation in the noise ε

$$\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$$

$$s^2 = \frac{1}{n-1} \sum_{i=1}^n (y_i - \bar{y})^2$$

$$\sum (y_i - \bar{y}) = 0$$

called `sigma` in `glance()` output ~~def.~~

- The RSE estimates the standard deviation of the error term ϵ (the RSS is divided by the appropriate degrees of freedom to give a "good" estimate of $\sigma = \sqrt{\text{Var}(\epsilon)}$)
- The RSE is a measure based on training data to evaluate the fit of the model (for inference) and needed to estimate the standard errors of $\hat{\beta}_j$ in classical theory!
- The RSE gives an idea of the size of the *irreducible* error, very similar to the RSS, small is good

$$y = \beta_0 + \beta_1 x + \epsilon$$

$$\epsilon \sim N(0, \sigma^2)$$

$$\sigma = \text{RSE}$$

Mean Squared Error: $MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$ *residual*

✓ *Training MSE*: if we use the original sample $y_1, \dots, y_n$ and their predicted values (formula above). It can be easily obtained from `.resid` column in `augment()` output

• *Testing MSE*: can be computed on new data y_{new} and their predicted values (formula above) to evaluate *out-of-sample* prediction performance

29

Poll questions and answers

30

Question 1: model.1

Which of the following codes would you use to estimate the model below:

$$\text{prot}_i = \beta_0 + \varepsilon_i$$

☒ A `dat_3genes %>% lm(prot ~ ., data = .) %>% tidy()`

☐ B `dat_3genes %>% group_by(gene) %>% do(model=tidy(lm(prot ~ 1, data = .))) %>% pull()`

☒ C `dat_3genes %>% lm(prot ~ 1, data = .) %>% tidy()`

☐ D `dat_3genes %>% lm(prot ~ mrna, data = .) %>% tidy()`

Answer: there is only one model without any covariate, the intercept is capture by 1 in the code. The notation `.` is used to include all other covariates in the data. In option B, a different model is fit to each gene using `group_by(gene)`

31

Question 2: model.2

Which of the following codes would you use to estimate the model below:

$$\text{prot}_i = \beta_0 + \beta_1 \times \text{mrna}_i + \varepsilon_i$$

☒ A `dat_3genes %>% lm(prot ~ mrna, data = .) %>% tidy()`

☐ B `dat_3genes %>% group_by(gene) %>% do(model=tidy(lm(prot ~ mrna, data = .))) %>% pull()`

☐ C `dat_3genes %>% lm(prot ~ mrna + gene, data = .) %>% tidy()`

Answer: there is only one model with `mrna` as the only covariate, a unique line. In option B, a different model is fit to each gene using `group_by(gene)`

32

Question 3: model.3

Which of the following codes would you use to estimate the model below:

$$\text{prot}_i = \beta_0 + \beta_1 \times \text{gene2}_i + \beta_2 \times \text{gene3}_i + \epsilon_i$$

A `dat_3genes %>% lm(prot ~ ., data = .) %>% tidy()`

B `dat_3genes %>% group_by(gene) %>% do(model=tidy(lm(prot ~ 1, data = .))) %>% pull()`

C `dat_3genes %>% lm(prot ~ 1, data = .) %>% tidy()`

✓ D `dat_3genes %>% lm(prot ~ gene, data = .) %>% tidy()`

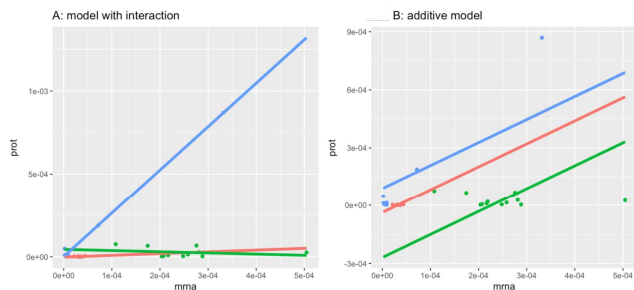
Answer: there is only one model without `mrna` and 2 dummy variables to identify each gene as the only covariate. This is not a line or plane since the covariates are not continuous. The covariates are used to distinguish the groups. Boxplots are good visualizations in this case. In option B, a different model is fit to each gene using `group_by(gene)`. Option C, does not make any distinction among genes.

33

Question 4: model.4

Which of the following plots represent the following model in:

$$\text{prot}_i = \beta_0 + \beta_1 \times \text{mrna}_i + \beta_2 \times \text{gene2}_i + \beta_3 \times \text{gene3}_i + \epsilon_i$$



34