

Logistic Regression

Estimation and Interpretations

In this lecture

- Logistic regression: definition
- Estimation
- Interpretation for different type of covariates
- Inference: confidence intervals and tests

Logistic Regression

[Scroll down to see full content](#)

Logistic regression is a powerful statistical method used for modelling a binary (and binomial) response variable.

For example, we can use a logistic regression to

1. compare the presence of bacteria between groups taking a new drug and a placebo, respectively.

- Response: *present or not present*

2. predict whether or not a customer will default on a loan given their income and demographic variables.

- Response: default or not default

© 2025 Gabriela Cohen Freue – Material Licensed under [CC By-SA 4.0](#)

3. know how GPA, ACT score, and number of AP classes taken are associated with the probability of getting accepted into a particular university.

- Response: accepted or not accepted

The response

[Scroll down to see full content](#)

To fit a logistic regression in R, the response needs to be a numerical variable (0 and 1) or a **factor**, with two levels (since R stores factors as integers).

For example, a numerical binary response Y_i that flags the successes (S) has the form:

$$Y_i = \begin{cases} 1 & \text{if the } i\text{th observation is S,} \\ 0 & \text{otherwise.} \end{cases}$$

In R, you can use the following code to convert a string response into a numerical one:

```
y_reponse <- if_else(y == "success", 1, 0)
```

Dataset

In this lecture, we'll use the [Titanic dataset](#), which contains survival and demographic information about the passengers.


- You can find details about the columns by typing `?titan`.

The `titanic` dataset has information about 891 passengers.

passenger_id	survived	passenger_class	name	sex	age	sibsp	parch	titanic
759	0	3	Theobald, Mr. Thomas Leonard	male	34	0	0	
523	0	3	Lahoud, Mr. Sarkis	male	NA	0	0	

Linear Regression: a range problem

Let's start by using a LR to explore if there is any association between **fare** and **survival**



A function with range (0, 1)

With a LR, fitted values may be larger than 1 or smaller than 0!


When the response is binary: $E[Y|X] = P[Y|X]$. We are modeling a **probability**!!

Let's use a **function of the linear component** with range (0, 1):

$$E[\text{survived}_i \mid \text{fare}_i] = P[\text{survived}_i \mid \text{fare}_i] = \frac{e^{\beta_0 + \beta_1 \times \text{fare}_i}}{1 + e^{\beta_0 + \beta_1 \times \text{fare}_i}}$$

Note that this ratio is always between 0 and 1!

SLR and Logistic Regression



The logit function

Using [some algebra](#) you can get an alternative formulation:

$$P [\text{survived}_i \mid \text{fare}_i] = p_i = \frac{e^{\beta_0 + \beta_1 \times \text{fare}_i}}{1 + e^{\beta_0 + \beta_1 \times \text{fare}_i}} \quad [\text{probability}]$$

can also be written as

$$\underbrace{\log\left(\frac{p_i}{1 - p_i}\right)}_{\text{the logit function}} = \beta_0 + \beta_1 \text{ fare}_i \quad [\text{log-odds}]$$

The Logistic Regression

We model the **logit of the probability** with a linear component!!

then all that we learned for MLR can be used for this function of the conditional expectation!!

Thus, this model is known as a **Logistic Regression**

$$\log\left(\frac{p_i}{1 - p_i}\right) = \beta_0 + \beta_1 X_{1i} + \beta_2 X_{2i} + \dots + \beta_p X_{pi}$$

Note that there is no error term! we are modeling a function of the conditional expectation

Logistic: a generalized linear model

The logistic regression belongs to a family of models called “generalized linear models”, aka GLMs.

In R, we can use the function `glm()` that estimates these models but we need to specify which model in the family we are interested in.

For logistic regression: `family = binomial`

In general, there is no closed-form to calculate the MLE. Thus, an iterative algorithm is needed (see # iteration is the `summary`)

The MLR also belongs to this family and you can also use `glm()` to estimate it.

It is the default family: `family = gaussian`

Note that `glm()` calculates maximum likelihood estimators (MLE) and `lm()` calculates least squares estimators.

For Normal random errors, both methods give the same estimates!

You can get more information about this function using `?glm`

Logit: log-odds

`glm()` models, by default, the **logit of the probabilities**, usually referred as the “canonical link” function

links the conditional expectation with a linear component

Note that the **logit of the probabilities** is the log of the odds of success (more details are given in the [math appendix](#))

$$\text{logit}(p) = \log \underbrace{\left(\frac{p}{1-p} \right)}_{\text{odds}} = \log(\text{odds})$$

Fitting a Logistic Regression

Case 1: one continuous covariate

[Scroll down to see full content](#)

$$\log\left(\frac{P(\text{survived}_i = 1 | \text{fare}_i)}{1 - P(\text{survived}_i = 1 | \text{fare}_i)}\right) = \beta_0 + \beta_1 \text{fare}_i$$

```
1 model_titanic_logistic <- glm(formula = survived ~ fare, data = titan,  
2                               family = 'binomial')  
3  
4 summary(model_titanic_logistic)
```

Call:

```
glm(formula = survived ~ fare, family = "binomial", data = titan)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.941330	0.095129	-9.895	< 2e-16 ***
fare	0.015197	0.002232	6.810	9.79e-12 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1186.7 on 890 degrees of freedom
Residual deviance: 1117.6 on 889 degrees of freedom
AIC: 1121.6

Interpretation: log-odds

The coefficients can be interpreted in terms of the log-odds (logit link):

- **Intercept** $\hat{\beta}_0 = -0.9413$: if the **fare** paid was 0 dollars, the log-odds of surviving are expected to be -0.9413.
- **Slope** $\hat{\beta}_1 = 0.0152$: an increase of 1 dollar in the **fare** paid is associated with an increase in the log-odds of surviving of 0.0152.

same as for MLR but with response “log-odds of success”!

Interpretation: odds

A more natural way of interpreting the coefficients of a logistic regression is based on **odds**

Using [basic algebra](#), the model can also be written as:

$$\frac{p_i}{1 - p_i} = e^{\beta_0 + \beta_1 \text{fare}_i} = e^{\beta_0} e^{\beta_1 \text{fare}_i} \quad [\text{odds}]$$

We need to exponentiate the coefficients to interpret them as changes in the **odds**. Find more details in the [math appendix](#)

Exponentiating the coefficients

```
1 #Coefficients for the odds
2 tidy(model_titanic_logistic, exponentiate = TRUE) %>%
3   mutate_if(is.numeric, round, 3) %>%
4   knitr::kable()
```

term	estimate	std.error	statistic	p.value
(Intercept)	0.390	0.095	-9.895	0
fare	1.015	0.002	6.810	0

Slope: $e^{\hat{\beta}_1} = e^{0.0152} = 1.015$: an increase of 1 dollar in the **fare** paid is associated with an increase of the odds of surviving **by a factor** of 1.015 or **an increase of 1.5% percent**

Case 2: one categorical variable

$$\log\left(\frac{P(\text{survived}_i = 1 \mid X_{\text{male},i})}{1 - P(\text{survived}_i = 1 \mid X_{\text{male},i})}\right) = \beta_0 + \beta_1 X_{\text{male},i}$$

where $X_{\text{male},i} = 1$ if passenger i is registered as a male and $X_{\text{male},i} = 0$ if registered as female.

Again, 2 logistic regressions in 1!!

$$\text{logit}(p_{\text{survival}}) = \beta_0 + \beta_1 \times X_{\text{male}} = \begin{cases} \beta_0 & \text{if } X_{\text{male}} = 0 \\ \beta_0 + \beta_1 & \text{if } X_{\text{male}} = 1 \end{cases}$$

Therefore, for the odds of surviving:

$$\frac{p_{\text{survival}}}{1 - p_{\text{survival}}} = \begin{cases} e^{\beta_0} & \text{if } X_{\text{male}} = 0 \\ e^{\beta_0 + \beta_1} & \text{if } X_{\text{male}} = 1 \end{cases}$$

find more details in the [math appendix](#)

in R

Scroll down to see full content

```
1 model_titanic_logistic_sex <- glm(survived ~ sex, data = titan, family = bi
2
3 tidy(model_titanic_logistic_sex) %>%
4   mutate_if(is.numeric, round, 3) %>%
5   knitr::kable()
```

term	estimate	std.error	statistic	p.value
(Intercept)	1.057	0.129	8.191	0
sexmale	-2.514	0.167	-15.036	0

```
tidy(model_titanic_logistic_sex, exponentiate = TRUE) %>%
  mutate_if(is.numeric, round, 3) %>%
  knitr::kable()
```

term	estimate	std.error	statistic	p.value
(Intercept)	2.877	0.129	8.191	0
sexmale	0.081	0.167	-15.036	0

- Reference level: female, note the dummy variable name sexmale

Interpretation (in terms of odds)

Scroll down to see full content

- **Intercept:** $e^{\hat{\beta}_0} = e^{1.057} = 2.8765$, a female passenger had a 2.8765 odds of surviving (i.e., the proportion of survivals relative to the proportion of deaths in the sample)
- **Slope:** $e^{\hat{\beta}_1} = e^{-2.514} = 0.08097$, the odds of surviving for a male passenger are **multiplied** by a factor of 0.08097; a decrease of 91.90% (note $(0.08097 - 1) \times 100 = -91.90$)

I

Odds ratio: note that $e^{\hat{\beta}_1} = 0.08097 = \frac{\text{odds}_{\text{male}}}{\text{odds}_{\text{female}}}$

22

- turning to **odds of failure**: the male passengers' **odds of dying** were $e^{2.514} = 12.3542$ **times** higher than of those of females

More details about these equivalences are given in the [math appendix](#)

Case 3: one categorical and one numerical covariate

Let's fit a logistic regression using **sex** and **fare** (omitting subscript i for simplicity)

The additive model

$$\log\left(\frac{P(\text{surv} = 1 \mid \text{fare}, X_{\text{male}})}{1 - P(\text{surv} = 1 \mid \text{fare}, X_{\text{male}})}\right) = \beta_0 + \beta_1 X_{\text{male}} + \beta_2 \times \text{fare}$$

Again, 2 logistic regressions in 1 with **equal** slope!


$$\text{logit}(p_{\text{survival}}) = \beta_0 + \beta_1 \times X_{\text{male}} + \beta_2 \times \text{fare}$$

$$\log\left(\frac{p_s}{1 - p_s}\right) = \begin{cases} \beta_0 + \beta_2 \times \text{fare}, & \text{if } X_{\text{male}} = 0 \\ \beta_0 + \beta_1 + \beta_2 \times \text{fare}, & \text{if } X_{\text{male}} = 1 \end{cases}$$

| same *slope* (β_2), different *intercepts* (β_0 vs ($\beta_0 + \beta_1$))!

For the odds of surviving:

$$\frac{p_{\text{survival}}}{1 - p_{\text{survival}}} = \begin{cases} e^{\beta_0} e^{\beta_2 \text{fare}}, & \text{if } X_{\text{male}} = 0 \\ e^{\beta_0 + \beta_1} e^{\beta_2 \text{fare}}, & \text{if } X_{\text{male}} = 1 \end{cases}$$



The curves won't look “parallel”, even if the linear components have the same slopes

$$p_i = \begin{cases} \frac{e^{\beta_0 + \beta_2 \times \text{fare}_i}}{1 + e^{\beta_0 + \beta_2 \times \text{fare}_i}} & \text{if } X_{\text{male}, i} = 0 \\ \frac{e^{\beta_0 + \beta_1 + \beta_2 \times \text{fare}_i}}{1 + e^{\beta_0 + \beta_1 + \beta_2 \times \text{fare}_i}} & \text{if } X_{\text{male}, i} = 1 \end{cases}$$

in R

Scroll down to see full content

```
1 model_titanic_logistic_additive <- glm(survived ~ sex + fare, data = titan,  
2                                     family = binomial)  
3  
4 tidy(model_titanic_logistic_additive) %>%  
5   mutate_if(is.numeric, round, 3) %>%  
6   knitr::kable()
```

term	estimate	std.error	statistic	p.value
(Intercept)	0.647	0.149	4.358	0
sexmale	-2.423	0.171	-14.208	0
fare	0.011	0.002	4.886	0

```
1 tidy(model_titanic_logistic_additive, exponentiate = TRUE) %>%  
2   mutate_if(is.numeric, round, 3) %>%  
3   knitr::kable()
```

term	estimate	std.error	statistic	p.value
(Intercept)	1.910	0.149	4.358	0
sexmale	0.089	0.171	-14.208	0
fare	1.011	0.002	4.886	0

Interpretation (for odds ratios)

[Scroll down to see full content](#)

Keeping the **fare** value constant, **at any value**:

- the men's expected **odds of surviving** are $e^{-2.42276} = 0.0887$ **times** that of a female
- equivalently, male passengers have only 8.87% of the odds of surviving a female passenger.
- or in other words, the men's expected **odds of dying** are $e^{2.42276} = 11.2769$ **times** that of a female

note that we can “keep **fare** constant” since this is an **additive** model. More details in the [math appendix](#)

Interpretation (cont.)

[Scroll down to see full content](#)

Keeping the **sex** constant, **for either male or female**:

- increasing the fare in 1 dollar is associated with an increase of the odds of surviving **by a factor of** 1.0113.
- increasing the fare in 1 dollar is associated with an increase of the odds of surviving of $(1.0113 - 1) \times 100\% = 1.13\%$.

Case 4: one categorical and one numerical covariate

Model with interaction

$$\log\left(\frac{p_s}{1 - p_s}\right) = \beta_0 + \beta_1 X_{\text{male}} + \beta_2 \times \text{fare} + \beta_3 \times X_{\text{male}} \times \text{fare}$$

Again, 2 logistic regressions in 1 with **different** slopes!


$$\text{logit}(p_s) = \beta_0 + \beta_1 \times X_{\text{male}} + \beta_2 \times \text{fare} + \beta_3 \times X_{\text{male}} \times \text{fare}$$

$$\log\left(\frac{p_s}{1 - p_s}\right) = \begin{cases} \beta_0 + \beta_2 \times \text{fare}, & \text{if } X_{\text{male}} = 0 \\ (\beta_0 + \beta_1) + (\beta_2 + \beta_3) \times \text{fare}, & \text{if } X_{\text{male}} = 1 \end{cases}$$

different *slope* (β_2 vs $(\beta_2 + \beta_3)$), different *intercepts* (β_0 vs $(\beta_0 + \beta_1)$)!

For the odds of surviving:

$$\frac{p_{\text{survival}}}{1 - p_{\text{survival}}} = \begin{cases} e^{\beta_0} e^{\beta_2 \text{fare}}, & \text{if } X_{\text{male}} = 0 \\ e^{(\beta_0 + \beta_1)} e^{(\beta_2 + \beta_3) \text{fare}}, & \text{if } X_{\text{male}} = 1 \end{cases}$$



$$p_i = \begin{cases} \frac{e^{\beta_0 + \beta_2 \times \text{fare}_i}}{1 + e^{\beta_0 + \beta_2 \times \text{fare}_i}} & \text{if } X_{\text{male}, i} = 0 \\ \frac{e^{(\beta_0 + \beta_1) + (\beta_2 + \beta_3) \times \text{fare}_i}}{1 + e^{(\beta_0 + \beta_1) + (\beta_2 + \beta_3) \times \text{fare}_i}} & \text{if } X_{\text{male}, i} = 1 \end{cases}$$

in R

Scroll down to see full content

```
1 model_titanic_logistic_int <- glm(survived ~ sex * fare, data = titan,  
2                                 family = binomial)  
3  
4 tidy(model_titanic_logistic_int) %>%  
5   mutate_if(is.numeric, round, 3) %>%  
6   knitr::kable()
```

term	estimate	std.error	statistic	p.value
(Intercept)	0.408	0.190	2.150	0.032
sexmale	-2.099	0.230	-9.116	0.000
fare	0.020	0.005	3.701	0.000
sexmale:fare	-0.012	0.006	-1.958	0.050

```
1 tidy(model_titanic_logistic_int, exponentiate = TRUE) %>%  
2   mutate_if(is.numeric, round, 3) %>%  
3   knitr::kable()
```

term	estimate	std.error	statistic	p.value
(Intercept)	1.504	0.190	2.150	0.032
sexmale	0.123	0.230	-9.116	0.000
fare	1.020	0.005	3.701	0.000
sexmale:fare	0.988	0.006	-1.958	0.050

Interpretation (for odds ratios)

[Scroll down to see full content](#)

In a **model with interactions**, the expected change in one variable **depends on** the value or levels the other variable

We **CAN'T** say “keeping the **fare** value constant” or “for either male or female”

- *for female passengers*, increasing the fare in 1 dollar is associated with an increase of the odds of surviving **by a factor of 1.020**.
- *for female passengers*, increasing the fare in 1 dollar is associated with an increase of the odds of surviving of 2%.

Interaction term

The interpretation of the interaction terms becomes a bit complicated.

Instead, you can calculate the two estimated curves and interpret them accordingly.

From the `tidy` table (with `exponentiate = FALSE` or default), we get

$$\hat{\beta}_0 = 0.408, \hat{\beta}_1 = -2.099, \hat{\beta}_2 = 0.020, \hat{\beta}_3 = -0.012$$

Then,

$$\hat{\beta}_0 = 0.408$$

$$\hat{\beta}_2 = 0.020$$

$$\hat{\beta}_0 + \hat{\beta}_1 = 0.408 - 2.099 = -1.691$$

$$\hat{\beta}_2 + \hat{\beta}_3 = 0.020 - 0.012 = 0.008$$

$$\frac{p_i}{1 - p_i} = \begin{cases} e^{0.408 + 0.020 \times \text{fare}_i} & \text{if } X_{\text{male}, i} = 0 \\ e^{-1.691 + 0.008 \times \text{fare}_i} & \text{if } X_{\text{male}, i} = 1 \end{cases}$$

- *for male passengers*, increasing the fare in 1 dollar is associated with an increase of the odds of surviving **by a factor of** $e^{0.008} = 1.008$.
- *for male passengers*, increasing the fare in 1 dollar is associated with an increase in the odds of surviving of 0.8%.

Using exponentiated coefficients

Note that you can also recover both curves using the exponentiated coefficients, from the `tidy` table, with `exponentiate = TRUE`.

$$e^{\hat{\beta}_0} = 1.504, e^{\hat{\beta}_1} = 0.123, e^{\hat{\beta}_2} = 1.020, e^{\hat{\beta}_3} = 0.988$$

We need to multiply (instead of summing):

$$e^{\hat{\beta}_0 + \hat{\beta}_1} = e^{-1.691} = 0.184$$

$$e^{\hat{\beta}_0} e^{\hat{\beta}_1} = 1.504 \times 0.123 = 0.185$$

$$e^{\hat{\beta}_2 + \hat{\beta}_3} = e^{0.008} = 1.008$$

$$e^{\hat{\beta}_2} e^{\hat{\beta}_3} = 1.020 \times 0.988 = 1.008$$

which are equivalent, up to rounding errors

Inference

Inference for GLMs, and logistic regression in particular, are based on theoretical results from the CLT.

if the sample size is large, the *Normal* distribution is used to obtain boundaries of confidence intervals and p -values

$$\frac{\hat{\beta}_j}{\hat{\text{SE}}(\hat{\beta}_j)} \sim N(0, 1)$$

This is called the **Wald's statistic**

The confidence intervals of $(1 - \alpha)\%$ for the coefficients of the log-odds model (canonical link) are

$$CI(\hat{\beta}_j, 1 - \alpha) = \hat{\beta}_j \pm z_{\alpha/2} \hat{SE}(\hat{\beta}_j)$$

You can also use the Wald's statistic to test $H_0 : \hat{\beta}_j = 0$ vs $H_1 : \hat{\beta}_j \neq 0$

But again, R will do all this for you!

in R

```
1 tidy(model_titanic_logistic_additive,  
2     conf.int = TRUE, conf.level = 0.9) %>%  
3   mutate_if(is.numeric, round, 3) %>%  
4   knitr::kable()
```

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	0.647	0.149	4.358	0	0.405	0.894
sexmale	-2.423	0.171	-14.208	0	-2.707	-2.145
fare	0.011	0.002	4.886	0	0.008	0.015

For **fare**: $p < 0.001$ indicating significant evidence to **reject the null hypothesis** that states that there is no association between **fare** and the log-odds of surviving, for either female or male passengers.

Inference for odds of success

Scroll down to see full content

You can also get confidence intervals for the exponentiated coefficients

```
1 tidy(model_titanic_logistic_additive,  
2     conf.int = TRUE, conf.level = 0.9, exponentiate = TRUE)%>%  
3 mutate_if(is.numeric, round, 3) %>%  
4 knitr::kable()
```

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	1.910	0.149	4.358	0	1.499	2.444
sexmale	0.089	0.171	-14.208	0	0.067	0.117
fare	1.011	0.002	4.886	0	1.008	1.015

We can be 90% confident that an increase of 1 dollar in **fare** is associated with a 0.8% to 1.5% increase in the odds of surviving for either male or female passengers.

since $(1.008 - 1) \times 100\% = 0.8\%$ and
 $(1.015 - 1) \times 100\% = 1.5\%$

Summary

- The (conditional) expectation of a binary response is the (conditional) *probability of success*.
- A LR can not be used to model the conditional expectation of a binary response since its range extends beyond the interval $(0, 1)$
- Instead, one can model a function of the conditional probability. A common choice in logistic regression is to use the *logit* function (logarithm of odds)

Summary (cont.)

[Scroll down to see full content](#)

The interpretation of the coefficients depends on the type of variables and the form of the model:

The raw coefficients (of the log-odds model) are interpreted as:

- log-odds of a reference group
- difference of log-odds of a treatment vs a control group
- changes in log-odds per unit change in the input

The exponentiated coefficients (of the odds model) are interpreted as:

- odds of a reference group
- odds ratio of a treatment vs a control group
- multiplicative changes in odds per unit change in the input

Summary (cont.)

- The estimated logistic model can be used to make inference using the Wald's statistic: confidence intervals and tests
- In the next lecture we'll look at
 - fitted values
 - residuals
 - diagnostics